

Contents

Project 4: Introduction to Deep Learning	1
Dependencies	1
Setup and Running Executables	1
Description	1
1. MNIST Tutorial	1
2. Experiment with Network Variations	2
Developing a Plan	2
Predict the Results	3
Results and Evaluation	3
3. Transfer Learning on Greek Letters	5
4. Heart Disease Prediction Using an ANN	7
The Networks	7
Computational Requirements and Performance	9
Extensions	9
Reflections	9
References	12

Project 4: Introduction to Deep Learning

Dependencies

This project has the following requirements:

```
pandas
matplotlib
plotly
numpy
torch
torchvision
ray[tune]
scikit-learn
```

Setup and Running Executables

All analysis in this project are conducted in either Jupyter Notebooks or using Python scripts. Tuning hyper-parameters and training models are done using Ray Tune, which can be installed using `pip install ray[tune]`.

Of note, some of the deep network training are done with M1 GPU acceleration, which may not be available on other machines. If you are using a different machine, you may need to change the device to `cpu` in the code.

The PDF report for this project is created from README.md using

```
pandoc README.md -o project4_report.pdf "-fmarkdown-implicit_figures -o"
--from=markdown -V geometry:margin=.8in --toc --highlight-style=espresso
```

Description

In this project, we begin using deep learning to solve a few classification problems with PyTorch. We explore different deep network architectures, other hyper-parameters, as well as using transfer learning to perform new tasks based on a pre-trained model.

1. MNIST Tutorial

In this section, we have implemented the tutorial for working with MNIST data and trained a model for classification handwritten letters.

This is the network architecture used:

```

Net(
  (conv1): Conv2d(1, 10, kernel_size=(5, 5), stride=(1, 1))
  (conv2): Conv2d(10, 20, kernel_size=(5, 5), stride=(1, 1))
  (conv2_drop): Dropout2d(p=0.5, inplace=False)
  (fc1): Linear(in_features=320, out_features=50, bias=True)
  (fc2): Linear(in_features=50, out_features=10, bias=True)
)
torch.Size([10, 1, 5, 5])
torch.Size([10])
torch.Size([20, 10, 5, 5])
torch.Size([20])
torch.Size([50, 320])
torch.Size([50])
torch.Size([10, 50])
torch.Size([10])

```

Here are some example outputs and predictions:




Figure 1: example output

We can visualize our training and test loss:

To further improve our network, we can also perform continued training, where we load model states that were saved during the first training. Here are the additional results:

2. Experiment with Network Variations

Developing a Plan

- Train with default config to get baseline results




Figure 2: training test loss

- Separate our data into training, validation and test sets
- Define search configs for hyper-parameters and set up pipelines
- Search through combinations with Ray Tune
- Evaluate and present the best model configs

Predict the Results

Currently, our accuracy is only 82%. This means that we need to train a more powerful network with more epochs. We predict that as we increase the number of convolution layers and the size of the fully connected layers, we will get better performance.

Results and Evaluation

Using `config`, we define the search space for our hyper-parameters. We can then use `tune.run` to search through the combinations of hyper-parameters. We can also use `tune.run` to train our model with the best hyper-parameters.

```
config = {
    "nn": tune.choice(["c1", "c2", "c3", "c4", "c5", "c6"]),
    "l1": tune.sample_from(lambda _: 2**np.random.randint(2, 9)),
    "l2": tune.sample_from(lambda _: 2**np.random.randint(2, 9)),
    "lr": tune.loguniform(1e-4, 1e-1),
    "batch_size": tune.choice([8, 16, 32, 64])
}
```

Using Ray Tune, we have trained a few different models with different hyper-parameters by random search through the configuration space. Using Tensorboard, we were able to visualize each model and upload the results to Tensorboard.dev.

Here is the result output from Ray Tune:




Figure 3: continued training test loss

Number of trials: 50/50 (50 TERMINATED)

Trial name	status	loc	batch_size	l1	l2	lr	nn
train_fmnist_66f43_00000	TERMINATED	127.0.0.1:39054	16	16	256	0.00206247	c3
train_fmnist_66f43_00001	TERMINATED	127.0.0.1:39054	8	32	256	0.0206503	c1
train_fmnist_66f43_00002	TERMINATED	127.0.0.1:39054	16	64	16	0.00197529	c5
train_fmnist_66f43_00003	TERMINATED	127.0.0.1:39054	16	4	128	0.00258987	c6
train_fmnist_66f43_00004	TERMINATED	127.0.0.1:39054	64	4	32	0.0674501	c3
train_fmnist_66f43_00005	TERMINATED	127.0.0.1:39054	8	32	32	0.00062723	c5
train_fmnist_66f43_00006	TERMINATED	127.0.0.1:39054	16	16	64	0.0011763	c3
train_fmnist_66f43_00007	TERMINATED	127.0.0.1:39054	16	128	8	0.0114439	c6
train_fmnist_66f43_00008	TERMINATED	127.0.0.1:39054	8	32	8	0.00626231	c3
train_fmnist_66f43_00009	TERMINATED	127.0.0.1:39054	8	8	32	0.000363216	c6
train_fmnist_66f43_00010	TERMINATED	127.0.0.1:39054	16	32	16	0.0612799	c7
train_fmnist_66f43_00011	TERMINATED	127.0.0.1:39054	32	128	128	0.00664651	c7
train_fmnist_66f43_00012	TERMINATED	127.0.0.1:39054	8	32	256	0.00189866	c2
train_fmnist_66f43_00013	TERMINATED	127.0.0.1:39054	64	32	16	0.000883322	c4
train_fmnist_66f43_00014	TERMINATED	127.0.0.1:39054	32	32	16	0.00215716	c3
train_fmnist_66f43_00015	TERMINATED	127.0.0.1:39054	64	64	128	0.0300275	c5
train_fmnist_66f43_00016	TERMINATED	127.0.0.1:39054	8	128	4	0.0139086	c6
train_fmnist_66f43_00017	TERMINATED	127.0.0.1:39054	8	128	32	0.000394455	c8
train_fmnist_66f43_00018	TERMINATED	127.0.0.1:39054	64	8	256	0.000668037	c5
train_fmnist_66f43_00019	TERMINATED	127.0.0.1:39054	32	32	16	0.000156712	c5
train_fmnist_66f43_00020	TERMINATED	127.0.0.1:39054	64	32	64	0.00976354	c7
train_fmnist_66f43_00021	TERMINATED	127.0.0.1:39054	64	256	4	0.0684767	c4
train_fmnist_66f43_00022	TERMINATED	127.0.0.1:39054	16	64	128	0.00262746	c6

train_fmniest_66f43_00023	TERMINATED	127.0.0.1:39054		16		4		64		0.00863927		c8
train_fmniest_66f43_00024	TERMINATED	127.0.0.1:39054		8		64		16		0.0361757		c4
train_fmniest_66f43_00025	TERMINATED	127.0.0.1:39054		64		128		256		0.00267729		c3
train_fmniest_66f43_00026	TERMINATED	127.0.0.1:39054		64		4		32		0.000110058		c5
train_fmniest_66f43_00027	TERMINATED	127.0.0.1:39054		8		8		256		0.00289519		c7
train_fmniest_66f43_00028	TERMINATED	127.0.0.1:39054		64		8		64		0.00215591		c7
train_fmniest_66f43_00029	TERMINATED	127.0.0.1:39054		8		64		128		0.0629795		c2
train_fmniest_66f43_00030	TERMINATED	127.0.0.1:39054		16		128		256		0.0196399		c3
train_fmniest_66f43_00031	TERMINATED	127.0.0.1:39054		8		4		16		0.0237702		c4
train_fmniest_66f43_00032	TERMINATED	127.0.0.1:39054		32		64		64		0.0125313		c4
train_fmniest_66f43_00033	TERMINATED	127.0.0.1:39054		16		128		256		0.0542981		c3
train_fmniest_66f43_00034	TERMINATED	127.0.0.1:39054		32		8		32		0.00667866		c7
train_fmniest_66f43_00035	TERMINATED	127.0.0.1:39054		64		4		4		0.0639286		c3
train_fmniest_66f43_00036	TERMINATED	127.0.0.1:39054		8		128		32		0.00275755		c2
train_fmniest_66f43_00037	TERMINATED	127.0.0.1:39054		64		128		4		0.000795191		c7
train_fmniest_66f43_00038	TERMINATED	127.0.0.1:39054		32		64		64		0.000101874		c6
train_fmniest_66f43_00039	TERMINATED	127.0.0.1:39054		16		8		128		0.0165327		c3
train_fmniest_66f43_00040	TERMINATED	127.0.0.1:39054		32		4		4		0.000576219		c5
train_fmniest_66f43_00041	TERMINATED	127.0.0.1:39054		64		32		64		0.00155077		c8
train_fmniest_66f43_00042	TERMINATED	127.0.0.1:39054		8		32		256		0.000518675		c1
train_fmniest_66f43_00043	TERMINATED	127.0.0.1:39054		32		64		8		0.0157333		c7
train_fmniest_66f43_00044	TERMINATED	127.0.0.1:39054		16		16		32		0.000467206		c1
train_fmniest_66f43_00045	TERMINATED	127.0.0.1:39054		64		256		256		0.00130333		c1
train_fmniest_66f43_00046	TERMINATED	127.0.0.1:39054		64		256		8		0.0199584		c5
train_fmniest_66f43_00047	TERMINATED	127.0.0.1:39054		16		8		8		0.00036141		c6
train_fmniest_66f43_00048	TERMINATED	127.0.0.1:39054		8		64		16		0.00333754		c4
train_fmniest_66f43_00049	TERMINATED	127.0.0.1:39054		16		8		8		0.000220345		c3
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----												

2023-03-16 20:55:13,511 INFO tune.py:798 -- Total run time: 6493.39 seconds (6493.35 seconds for the tuning)
Best trial config: {'nn': 'c5', 'l1': 256, 'l2': 8, 'lr': 0.019958412992905763, 'batch_size': 64}
Best trial final validation loss: 0.3919870887506515
Best trial final validation accuracy: 0.85825

Best trial test set accuracy: 0.8536

Full results can be found on Tensorboard.dev at <https://tensorboard.dev/experiment/WgVGNysDQBa5PrkOY8Tn2w/>.

From this experiment, we can see that the best model is a 5-layer network with 3 convolution layers and 2 fully connected layers with 256 units in the first layer, 8 units in the second layer, and a learning rate of 0.0199. The best model achieves a test accuracy of 0.8536, which is slightly better than our baseline model. One can imagine that as we increase the number of convolution layers we may be able to achieve even better results.

3. Transfer Learning on Greek Letters

Next, we will use transfer learning to train a model to classify Greek letters. We will use the pre-trained model from part 1 and modify the last layer to classify the Greek letters.

First, we can visualize the greek letters after transformation:

Here is the structure of the modified network:

```
Net(
  (conv1): Conv2d(1, 10, kernel_size=(5, 5), stride=(1, 1))
  (conv2): Conv2d(10, 20, kernel_size=(5, 5), stride=(1, 1))
  (conv2_drop): Dropout2d(p=0.5, inplace=False)
  (fc1): Linear(in_features=320, out_features=50, bias=True)
```




Figure 4: greek letters

```

(fc2): Linear(in_features=50, out_features=3, bias=True)
)
torch.Size([10, 1, 5, 5])
torch.Size([10])
torch.Size([20, 10, 5, 5])
torch.Size([20])
torch.Size([50, 320])
torch.Size([50])
torch.Size([3, 50])
torch.Size([3])

```

Here are some example outputs and predictions:




Figure 5: example output

Now, we can visualize our training loss: it took about 200 epochs to reach an almost perfect accuracy.

4. Heart Disease Prediction Using an ANN

Now that we have had some practice with PyTorch, we will use it to predict heart disease using the heart disease data from the last project where we implement an ANN with linear layers.

The Networks

For making predictions on the heart disease data, we will experiment with two different network architectures and compare their performance. The first network has one hidden layer with 30 nodes and the second network has two hidden layers with 5 and 10 nodes each. We will use ReLU activation functions for both networks. We will also use the Adam optimizer with a learning rate of 0.01 and train both networks for 200 epochs.

LinearNN(

Greek Letter Training Loss



Figure 6: training loss


```

(fc1): Linear(in_features=12, out_features=30, bias=True)
(activation): ReLU()
(fc2): Linear(in_features=30, out_features=2, bias=True)
(softmax): Softmax(dim=None)
)
torch.Size([30, 12])
torch.Size([30])
torch.Size([2, 30])
torch.Size([2])

LinearNN(
(fc1): Linear(in_features=12, out_features=5, bias=True)
(activation): ReLU()
(fc2): Linear(in_features=5, out_features=10, bias=True)
(fc3): Linear(in_features=10, out_features=2, bias=True)
(softmax): Softmax(dim=None)
)
torch.Size([5, 12])
torch.Size([5])
torch.Size([10, 5])
torch.Size([10])
torch.Size([2, 10])
torch.Size([2])

```

Computational Requirements and Performance

Comparing the number of parameters in each network, we can see that the first network has 452 parameters while the second network has 147 parameters. This means that the first network is more powerful and will require more computational resources to train. However, the second network will be faster to train and will require less memory.

Interestingly, the first network has a higher accuracy than the second network. This is likely because the first network has more parameters and is therefore more powerful. The training time for the first network on the author's machine is 10.8s vs. 19.4s for the second network - we suspect that this is due to the difference in the optimizer SGD vs. Adam.

In comparing the efficiency of the training process, it is clear that the second network is more efficient.

Model 1

```

Number of parameters: 452
Training time: 10.8s
Accuracy: 0.82

```

Model 2

```

Number of parameters: 147
Training time: 19.4s
Accuracy: 0.79

```

We can also see that 200 epoch is probably too many for this dataset. Interestingly, ADAM optimizer was able to reduce the training loss much faster than the SGD optimizer.

Extensions

1. Transfer learning tutorial is implemented in `project4_extension.ipynb`.
2. Explored additional tuning methods and hyper-parameters in part 2.

Reflections

- PyTorch is a powerful tool for building neural networks. It seems easy to use while still being flexible enough to build complex models.

Training Loss and Test Loss




Figure 7: Heart Disease model1 training test loss

Training Loss and Test Loss

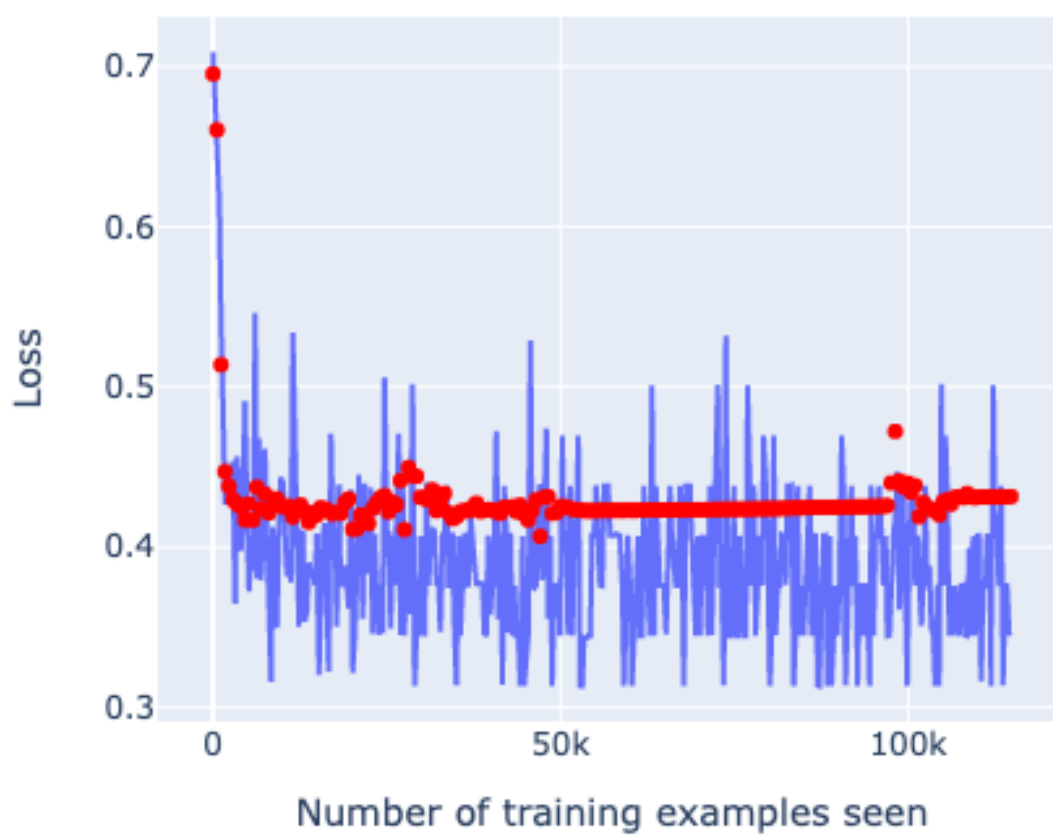


Figure 8: Heart Disease model2 training test loss

- It takes time to understand the different functions in PyTorch and how to use them. The tutorials are very helpful for this.
- Without very much experience, it was challenging to choose the right hyper-parameters for the models. Though experimenting with Ray Tune was very helpful automating the search through the configuration space.
- Transfer learning seems extremely useful for building a network for more sophisticated tasks with less data and reduced training time.
- Using M1 GPU acceleration significantly reduced the training time for some of the tasks while it caused the kernel to crash in others.
- Tensorboard is great for visualizing tuning results with the ability to upload it to Tensorboard.dev.

References

1. <https://towardsdatascience.com/hyperparameter-tuning-of-neural-networks-with-optuna-and-pytorch-22e179efc837>
2. <https://machinelearningmastery.com/choose-an-activation-function-for-deep-learning/>
3. <https://dingyan89.medium.com/calculating-parameters-of-convolutional-and-fully-connected-layers-with-keras-186590df36c6#:~:text=The%20kernel%20size%20of%20max,5%2C5%2C16>
4. <https://towardsdatascience.com/hyperparameter-tuning-of-neural-networks-with-optuna-and-pytorch-22e179efc837>
5. <https://shashikachamod4u.medium.com/excel-csv-to-pytorch-dataset-def496b6bcc1>
6. https://pytorch.org/tutorials/beginner/hyperparameter_tuning_tutorial.html
7. https://pytorch.org/tutorials/beginner/basics/data_tutorial.html
8. <https://machinelearningmastery.com/adam-optimization-algorithm-for-deep-learning/>
9. <https://tensorboard.dev/#add-name-description-to-experiment>
10. https://pytorch.org/tutorials/beginner/basics/data_tutorial.html
11. <https://towardsdatascience.com/installing-pytorch-on-apple-m1-chip-with-gpu-acceleration-3351dc44d67c>
12. <https://docs.ray.io/en/latest/ray-overview/getting-started.html>
13. <https://machinelearningmastery.com/adam-optimization-algorithm-for-deep-learning/>